

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder, OneHotEncoder,
StandardScaler, normalize
from sklearn.compose import ColumnTransformer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix
import tensorflow as tf

from google.colab import drive
drive.mount('/content/drive')

Drive already mounted at /content/drive; to attempt to forcibly
remount, call drive.mount("/content/drive", force_remount=True).

path =
"/content/drive/MyDrive/customer_churn/dataset/customer_churn_large_da
taset.xlsx"

df = pd.read_excel(path)

# Converting to Pandas DataFrame
df = pd.DataFrame(df)
df.head()

```

	CustomerID	Name	Age	Gender	Location	\
0	1	Customer_1	63	Male	Los Angeles	
1	2	Customer_2	62	Female	New York	
2	3	Customer_3	24	Female	Los Angeles	
3	4	Customer_4	36	Female	Miami	
4	5	Customer_5	46	Female	Miami	

	Subscription_Length_Months	Monthly_Bill	Total_Usage_GB	Churn
0	17	73.36	236	0
1	1	48.76	172	0
2	5	85.47	460	0
3	3	97.94	297	1
4	19	58.14	266	0

```

df.shape
(100000, 9)

df.dtypes

```

```

CustomerID      int64
Name            object
Age            int64
Gender          object
Location        object
Subscription_Length_Months  int64
Monthly_Bill    float64
Total_Usage_GB  int64
Churn           int64
dtype: object

```

```
df.describe(include='all')
```

	CustomerID	Name	Age	Gender	Location	\
count	100000.000000	100000	100000.000000	100000	100000	
unique	NaN	100000	NaN	2	5	
top	NaN	Customer_1	NaN	Female	Houston	
freq	NaN	1	NaN	50216	20157	
mean	50000.500000	NaN	44.027020	NaN	NaN	
std	28867.657797	NaN	15.280283	NaN	NaN	
min	1.000000	NaN	18.000000	NaN	NaN	
25%	25000.750000	NaN	31.000000	NaN	NaN	
50%	50000.500000	NaN	44.000000	NaN	NaN	
75%	75000.250000	NaN	57.000000	NaN	NaN	
max	100000.000000	NaN	70.000000	NaN	NaN	

	Subscription_Length_Months	Monthly_Bill	Total_Usage_GB	\
count	100000.000000	100000.000000	100000.000000	
unique	NaN	NaN	NaN	
top	NaN	NaN	NaN	
freq	NaN	NaN	NaN	
mean	12.490100	65.053197	274.393650	
std	6.926461	20.230696	130.463063	
min	1.000000	30.000000	50.000000	
25%	6.000000	47.540000	161.000000	
50%	12.000000	65.010000	274.000000	
75%	19.000000	82.640000	387.000000	
max	24.000000	100.000000	500.000000	

	Churn
count	100000.000000
unique	NaN
top	NaN
freq	NaN
mean	0.497790
std	0.499998
min	0.000000
25%	0.000000
50%	0.000000

75%	1.000000
max	1.000000

**Count** for each column is 100,000 which is the total number of data points

### Location

- Number of locations = 5
- Most common occurring location is **Houston** with 20157 data points

### Gender

- Majority of the customers are **Female**
- Total number of females = 50216

### Age

- Minimum age = 18
- Average age = 44
- Maximum Age = 70

### Subscription Duration

- Minimum = 1 month
- Average = 1 year
- Maximum = 2 years

=> There are no subscriptions that are for more than 2 years

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   CustomerID                           100000 non-null  int64
1   Name                                  100000 non-null  object
2   Age                                   100000 non-null  int64
3   Gender                               100000 non-null  object
4   Location                             100000 non-null  object
5   Subscription_Length_Months           100000 non-null  int64
6   Monthly_Bill                         100000 non-null  float64
7   Total_Usage_GB                       100000 non-null  int64
8   Churn                                100000 non-null  int64
dtypes: float64(1), int64(5), object(3)
memory usage: 6.9+ MB
```

=> There are no null values present in the dataset

##This can easily be confirmed using the isnull() method as well as visualized using a bar plot

```
df.isnull().any()


CustomerID      False
Name            False
Age            False
Gender          False
Location        False
Subscription_Length_Months  False
Monthly_Bill    False
Total_Usage_GB  False
Churn           False
dtype: bool
```

=> This indicates that none of the features contains any null values.

Visualization for the same can be seen below:

```
df.count().plot.bar()

<Axes: >
```



## Outlier Detection

### Visualization of outliers using Distribution

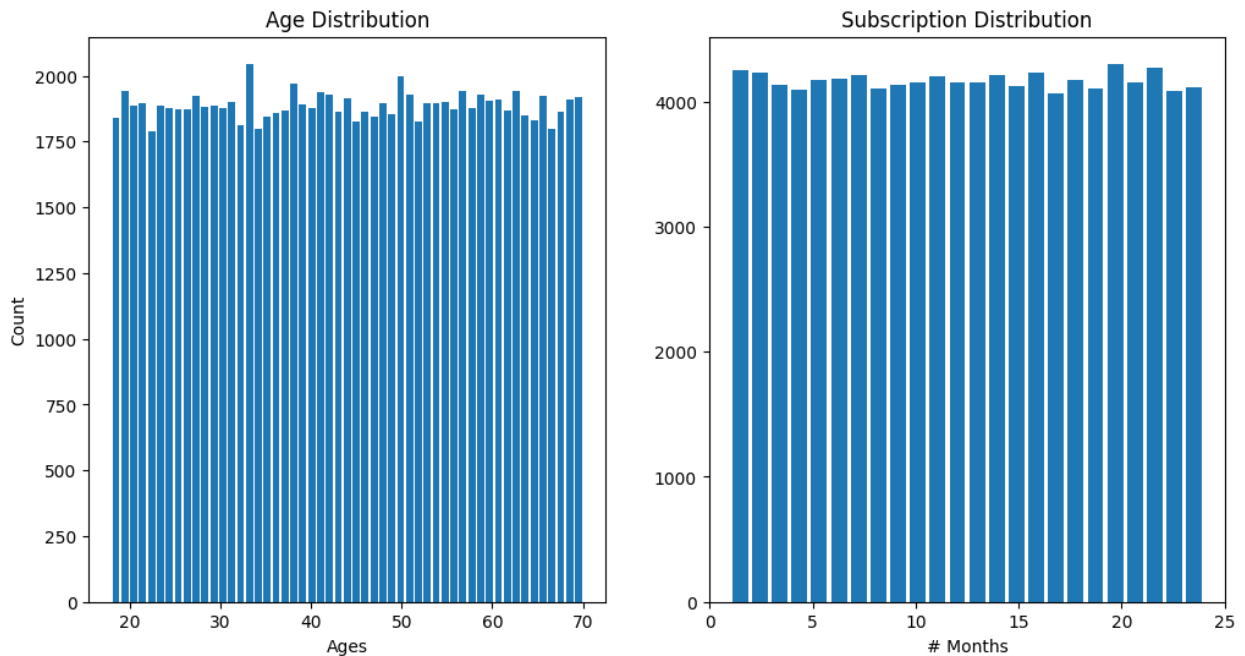
```
# Creating a 1x2 grid of subplots
fig, axs = plt.subplots(1, 2, figsize=(12, 6))

# Histogram of Age
axs[0].hist(df.Age, bins=(df.Age.max()-df.Age.min()+1, rwidth=0.8)
axs[0].set_title("Age Distribution")

# Histogram of subscription months
axs[1].hist(df.Subscription_Length_Months,
```

```
bins=(df.Subscription_Length_Months.max()-
df.Subscription_Length_Months.min()+1,rwidth=0.8)
axs[1].set_title("Subscription Distribution")

axs[0].set_ylabel('Count')
axs[0].set_xlabel('Ages')
axs[1].set_xlabel('# Months')
plt.show()
```



## Age Distribution


Age histogram exhibits uniform distribution along with the repeated peaks. Same pattern is followed throughout which rules out any possibility of outliers

## Subscription Distribution

- The even distribution of subscriptions suggests that people are subscribing uniformly across various time intervals.
- It that there are no particular periods that attract significantly more or fewer subscriptions than others

```
# Histogram of total usage
plt.hist(df.Total_Usage_GB, bins=20, rwidth=0.8)
```

```
plt.title("Total Usage Distribution")
plt.show()
```



- There is a uniform distribution of usage across various data packages which indicates that each package is utilized by similar proportion of the total population
- This could imply that there is no strong preference or bias towards any particular package
- It is safe to say that there are no outliers

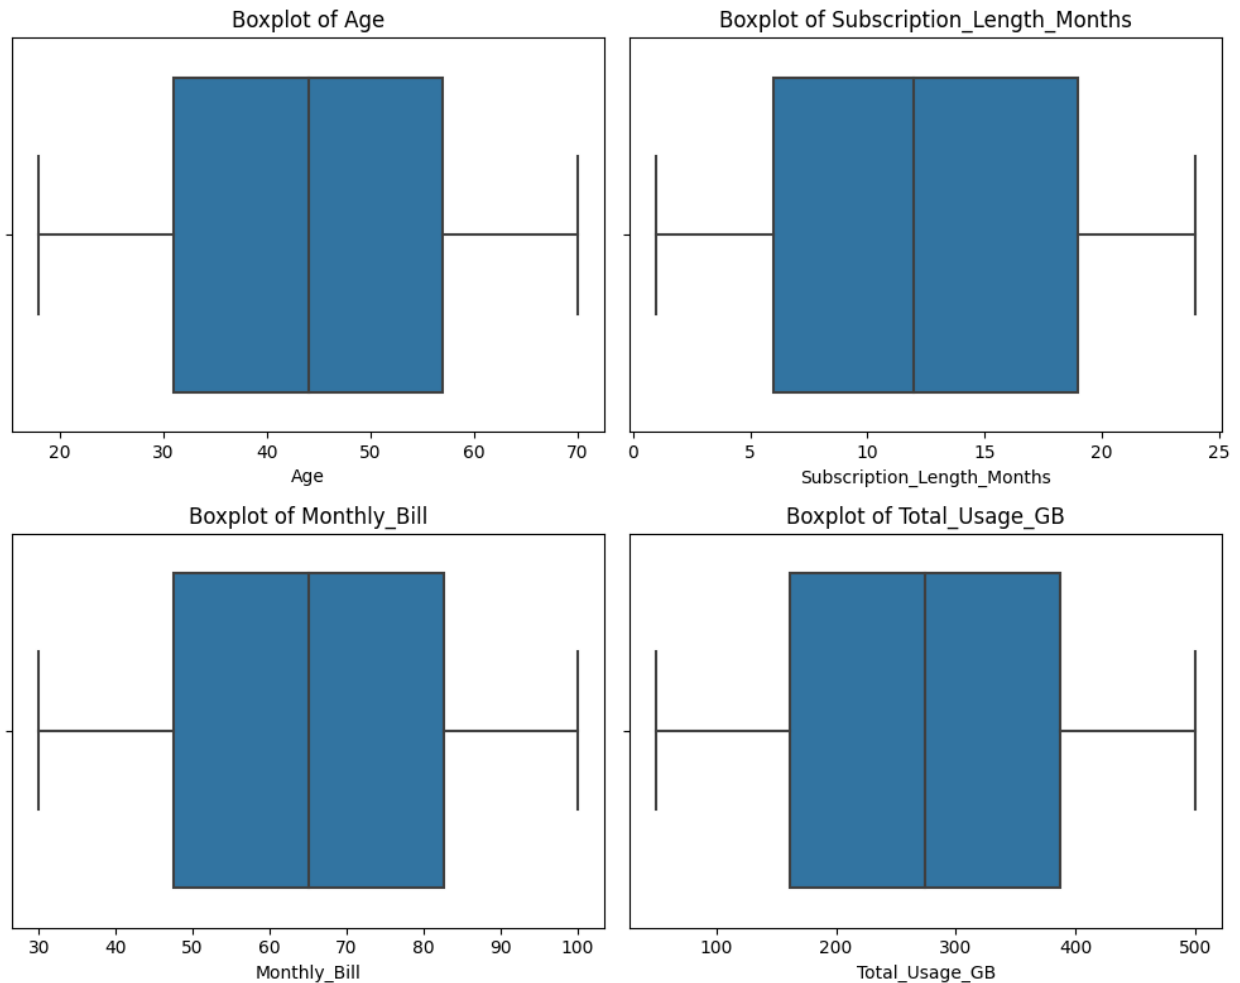
```
import seaborn as sns

# columns to analyze outliers
numerical_columns = ['Age', 'Subscription_Length_Months',
                    'Monthly_Bill', 'Total_Usage_GB']

# Plotting boxplots for the numerical columns
plt.figure(figsize=(10, 8))
for idx, col in enumerate(numerical_columns, 1):
    plt.subplot(2, 2, idx)
```

```
sns.boxplot(x=df[col])
plt.title(f'Boxplot of {col}')

plt.tight_layout()
plt.show()
```



## Outlier Detection using Z-Score

Considering threshold value to be 3. If the z-score  $> 3$ , this implies existence of outliers

```
# Age
z_score_age = (df.Age - df.Age.mean())/df.Age.std()
z_score_age
```

0	1.241664
1	1.176220



```

2      -1.310645
3      -0.525319
4       0.129119
...
99995   -0.721650
99996    1.176220
99997    1.307108
99998    0.456338
99999   -1.114313
Name: Age, Length: 100000, dtype: float64

df.Age[z_score_age >=3].count()

0

# Monthly Bill
z_score_bill = (df.Monthly_Bill -
df.Monthly_Bill.mean())/df.Monthly_Bill.std()
z_score_bill

0      0.410604
1     -0.805370
2      1.009199
3      1.625589
4     -0.341718
...
99995   -0.490502
99996   -0.168219
99997    1.535133
99998   -0.781149
99999    0.569274
Name: Monthly_Bill, Length: 100000, dtype: float64

df.Monthly_Bill[z_score_bill >=3].count()

0

# Total_Usage_GB
# Monthly Bill
z_score_usage = (df.Total_Usage_GB -
df.Total_Usage_GB.mean())/df.Total_Usage_GB.std()
z_score_usage

0     -0.294288
1     -0.784848
2      1.422674
3      0.173278
4     -0.064337
...
99995   -0.370938
99996    0.587188

```

```

99997    -0.179312
99998     1.223383
99999    -0.777183
Name: Total_Usage_GB, Length: 100000, dtype: float64

df.Total_Usage_GB[z_score_bill >=3].count()

0

```

The Z-score test implies that there are no existence of outliers in the features - Age, Monthly Bill and Total Usage

```

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 9 columns):
 #   Column                                Non-Null Count  Dtype  
---  -
 0   CustomerID                          100000 non-null  int64  
 1   Name                                100000 non-null  object  
 2   Age                                  100000 non-null  int64  
 3   Gender                              100000 non-null  object  
 4   Location                            100000 non-null  object  
 5   Subscription_Length_Months          100000 non-null  int64  
 6   Monthly_Bill                        100000 non-null  float64 
 7   Total_Usage_GB                      100000 non-null  int64  
 8   Churn                               100000 non-null  int64  
dtypes: float64(1), int64(5), object(3)
memory usage: 6.9+ MB

```

## Encoding Categorical Data

```

df.head()

```

	CustomerID	Name	Age	Gender	Location	\
0	1	Customer_1	63	Male	Los Angeles	
1	2	Customer_2	62	Female	New York	
2	3	Customer_3	24	Female	Los Angeles	
3	4	Customer_4	36	Female	Miami	
4	5	Customer_5	46	Female	Miami	

	Subscription_Length_Months	Monthly_Bill	Total_Usage_GB	Churn
0	17	73.36	236	0
1	1	48.76	172	0
2	5	85.47	460	0
3	3	97.94	297	1
4	19	58.14	266	0

There are two categorical features that need to be encoded, namely, Gender and Location

Extracting unique values from each feature

```
df['Gender'].unique()  
array(['Male', 'Female'], dtype=object)
```

=> Two unique values:

- Male
- Female

```
df['Location'].unique()  
array(['Los Angeles', 'New York', 'Miami', 'Chicago', 'Houston'],  
      dtype=object)
```

=> Five unique values:

- Los Angeles
- New York
- Miami
- Chicago
- Houston

df					
	CustomerID	Name	Age	Gender	Location \
0	1	Customer_1	63	Male	Los Angeles
1	2	Customer_2	62	Female	New York
2	3	Customer_3	24	Female	Los Angeles
3	4	Customer_4	36	Female	Miami
4	5	Customer_5	46	Female	Miami
...	...	...	...	...	...
99995	99996	Customer_99996	33	Male	Houston
99996	99997	Customer_99997	62	Female	New York
99997	99998	Customer_99998	64	Male	Chicago
99998	99999	Customer_99999	51	Female	New York
99999	100000	Customer_100000	27	Female	Los Angeles
	Subscription_Length_Months		Monthly_Bill	Total_Usage_GB	Churn
0	17		73.36	236	0
1	1		48.76	172	0
2	5		85.47	460	0
3	3		97.94	297	1

4	19	58.14	266	0
...	...	...	...	...
99995	23	55.13	226	1
99996	19	61.65	351	0
99997	17	96.11	251	1
99998	20	49.25	434	1
99999	19	76.57	173	1

```
[100000 rows x 9 columns]
```

## Applying Label-Encoding on Gender Feature

using sklearn

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 100000 entries, 0 to 99999
```

Data columns (total 9 columns):

#	Column	Non-Null Count		Dtype
0	CustomerID	100000	non-null	int64
1	Name	100000	non-null	object
2	Age	100000	non-null	int64
3	Gender	100000	non-null	object
4	Location	100000	non-null	object
5	Subscription_Length_Months	100000	non-null	int64
6	Monthly_Bill	100000	non-null	float64
7	Total_Usage_GB	100000	non-null	int64
8	Churn	100000	non-null	int64

```
dtypes: float64(1), int64(5), object(3)
```

```
memory usage: 6.9+ MB
```

## # Label Encoding

```
le = LabelEncoder()
```

```
df['Gender'] = le.fit_transform(df['Gender'])
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 100000 entries, 0 to 99999
```

Data columns (total 9 columns):

#	Column	Non-Null Count	Dtype
---	--------	----------------	-------

```

---
0  CustomerID      100000 non-null int64
1  Name           100000 non-null object
2  Age            100000 non-null int64
3  Gender          100000 non-null int64
4  Location        100000 non-null object
5  Subscription_Length_Months 100000 non-null int64
6  Monthly_Bill    100000 non-null float64
7  Total_Usage_GB  100000 non-null int64
8  Churn           100000 non-null int64
dtypes: float64(1), int64(6), object(2)
memory usage: 6.9+ MB

```

## Applying One Hot Encoding on Locations Feature

```

# Saving column names
dummy_cols = pd.get_dummies(df['Location']).columns
dummy_cols = list(dummy_cols)
print(dummy_cols)
cols = dummy_cols + ['CustomerId', 'Name', 'Age', 'Gender'] +
['Subscription_Length_Months', 'Monthly_Bill', 'Total_Usage_GB', 'Churn']
cols

```

```
['Chicago', 'Houston', 'Los Angeles', 'Miami', 'New York']
```

```

['Chicago',
 'Houston',
 'Los Angeles',
 'Miami',
 'New York',
 'CustomerId',
 'Name',
 'Age',
 'Gender',
 'Subscription_Length_Months',
 'Monthly_Bill',
 'Total_Usage_GB',
 'Churn']

```

df

	CustomerId	Name	Age	Gender	Location	\
0	1	Customer_1	63	1	Los Angeles	
1	2	Customer_2	62	0	New York	
2	3	Customer_3	24	0	Los Angeles	
3	4	Customer_4	36	0	Miami	
4	5	Customer_5	46	0	Miami	
...	...	...	...	...	...	
99995	99996	Customer_99996	33	1	Houston	
99996	99997	Customer_99997	62	0	New York	

99997	99998	Customer_99998	64	1	Chicago
99998	99999	Customer_99999	51	0	New York
99999	100000	Customer_100000	27	0	Los Angeles
		Subscription_Length_Months	Monthly_Bill	Total_Usage_GB	Churn
0		17	73.36	236	0
1		1	48.76	172	0
2		5	85.47	460	0
3		3	97.94	297	1
4		19	58.14	266	0
...		...	...	...	...
99995		23	55.13	226	1
99996		19	61.65	351	0
99997		17	96.11	251	1
99998		20	49.25	434	1
99999		19	76.57	173	1

[100000 rows x 9 columns]

```
ct = ColumnTransformer(transformers = [('encoder', OneHotEncoder(),
[4])], remainder='passthrough')
df = pd.DataFrame(ct.fit_transform(df))
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 13 columns):
#   Column  Non-Null Count  Dtype
---  -
0    0      100000 non-null  object
1    1      100000 non-null  object
2    2      100000 non-null  object
3    3      100000 non-null  object
4    4      100000 non-null  object
5    5      100000 non-null  object
6    6      100000 non-null  object
7    7      100000 non-null  object
8    8      100000 non-null  object
```

```
9      9      100000 non-null object
10     10      100000 non-null object
11     11      100000 non-null object
12     12      100000 non-null object
```

```
dtypes: object(13)
```

```
memory usage: 9.9+ MB
```

```
df.columns = cols
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 100000 entries, 0 to 99999
```

```
Data columns (total 13 columns):
```

#	Column	Non-Null Count	Dtype
0	Chicago	100000 non-null	object
1	Houston	100000 non-null	object
2	Los Angeles	100000 non-null	object
3	Miami	100000 non-null	object
4	New York	100000 non-null	object
5	CustomerId	100000 non-null	object
6	Name	100000 non-null	object
7	Age	100000 non-null	object
8	Gender	100000 non-null	object
9	Subscription_Length_Months	100000 non-null	object
10	Monthly_Bill	100000 non-null	object
11	Total_Usage_GB	100000 non-null	object
12	Churn	100000 non-null	object

```
dtypes: object(13)
```

```
memory usage: 9.9+ MB
```

```
df[cols] = df[cols].apply(pd.to_numeric, errors='coerce')
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 100000 entries, 0 to 99999
```

```
Data columns (total 13 columns):
```

#	Column	Non-Null Count	Dtype
0	Chicago	100000 non-null	float64
1	Houston	100000 non-null	float64
2	Los Angeles	100000 non-null	float64
3	Miami	100000 non-null	float64
4	New York	100000 non-null	float64
5	CustomerId	100000 non-null	int64
6	Name	0 non-null	float64
7	Age	100000 non-null	int64
8	Gender	100000 non-null	int64
9	Subscription_Length_Months	100000 non-null	int64
10	Monthly_Bill	100000 non-null	float64

```

11 Total_Usage_GB          100000 non-null int64
12 Churn                   100000 non-null int64
dtypes: float64(7), int64(6)
memory usage: 9.9 MB

```

## Train Test Splitting

```

y = np.array(df['Churn'])
y
df = df.drop(columns=['Churn'], axis = 1)
df

```

	Chicago	Houston	Los Angeles	Miami	New York	CustomerId
Name Age \						
0	0.0	0.0	1.0	0.0	0.0	1
NaN 63						
1	0.0	0.0	0.0	0.0	1.0	2
NaN 62						
2	0.0	0.0	1.0	0.0	0.0	3
NaN 24						
3	0.0	0.0	0.0	1.0	0.0	4
NaN 36						
4	0.0	0.0	0.0	1.0	0.0	5
NaN 46						
...	...	...	...	...	...	...
99995	0.0	1.0	0.0	0.0	0.0	99996
NaN 33						
99996	0.0	0.0	0.0	0.0	1.0	99997
NaN 62						
99997	1.0	0.0	0.0	0.0	0.0	99998
NaN 64						
99998	0.0	0.0	0.0	0.0	1.0	99999
NaN 51						
99999	0.0	0.0	1.0	0.0	0.0	100000
NaN 27						

	Gender	Subscription_Length_Months	Monthly_Bill
Total_Usage_GB			
0	1	17	73.36
236			
1	0	1	48.76
172			
2	0	5	85.47
460			
3	0	3	97.94
297			



```

4          0          19          58.14
266
...          ...          ...          ...
.
99995      1          23          55.13
226
99996      0          19          61.65
351
99997      1          17          96.11
251
99998      0          20          49.25
434
99999      0          19          76.57
173

```

```
[100000 rows x 12 columns]
```

```
# drop CustomerId column
```

```
df = df.drop(columns=['CustomerId', 'Name'],axis=1)
```

```
X = df
```

```
X
```

```

          Chicago  Houston  Los Angeles  Miami  New York  Age  Gender  \
0          0.0      0.0          1.0    0.0          0.0   63      1
1          0.0      0.0          0.0    0.0          1.0   62      0
2          0.0      0.0          1.0    0.0          0.0   24      0
3          0.0      0.0          0.0    1.0          0.0   36      0
4          0.0      0.0          0.0    1.0          0.0   46      0
...          ...      ...          ...    ...          ...   ...   ...
99995      0.0      1.0          0.0    0.0          0.0   33      1
99996      0.0      0.0          0.0    0.0          1.0   62      0
99997      1.0      0.0          0.0    0.0          0.0   64      1
99998      0.0      0.0          0.0    0.0          1.0   51      0
99999      0.0      0.0          1.0    0.0          0.0   27      0

```

```

          Subscription_Length_Months  Monthly_Bill  Total_Usage_GB
0                                17          73.36          236
1                                 1          48.76          172
2                                 5          85.47          460
3                                 3          97.94          297
4                                19          58.14          266
...          ...          ...          ...
99995                        23          55.13          226
99996                        19          61.65          351
99997                        17          96.11          251
99998                        20          49.25          434
99999                        19          76.57          173

```

```
[100000 rows x 10 columns]
```

```
# splitting the dataset
```

```
X_train,X_test,y_train,y_test =  
train_test_split(X,y,test_size=0.2,random_state=42)
```

```
X_train
```

	Chicago	Houston	Los Angeles	Miami	New York	Age	Gender	\
75220	0.0	0.0	0.0	0.0	1.0	54	0	
48955	0.0	0.0	0.0	0.0	1.0	28	1	
44966	1.0	0.0	0.0	0.0	0.0	57	1	
13568	0.0	1.0	0.0	0.0	0.0	19	1	
92727	0.0	0.0	0.0	1.0	0.0	56	0	
...	...	...	...	...	...	...	...	
6265	0.0	0.0	0.0	1.0	0.0	35	1	
54886	1.0	0.0	0.0	0.0	0.0	56	1	
76820	0.0	1.0	0.0	0.0	0.0	69	1	
860	1.0	0.0	0.0	0.0	0.0	55	1	
15795	0.0	0.0	1.0	0.0	0.0	26	0	

	Subscription_Length_Months	Monthly_Bill	Total_Usage_GB
75220	5	84.50	205
48955	24	82.06	239
44966	12	52.29	62
13568	19	32.57	173
92727	8	33.52	314
...	...	...	...
6265	21	67.33	235
54886	13	85.40	347
76820	2	76.24	321
860	12	89.19	315
15795	17	70.41	335

```
[80000 rows x 10 columns]
```

## Feature Scaling

### Standardization (z-score) - Normalization

```
X_train.describe()
```

	Chicago	Houston	Los Angeles	Miami	New York
count	80000.000000	80000.000000	80000.000000	80000.000000	80000.000000
mean	0.198313	0.201387	0.200175	0.201513	0.198612
std	0.398731	0.401039	0.400134	0.401132	0.398958

min	0.000000	0.000000	0.000000	0.000000
0.000000				
25%	0.000000	0.000000	0.000000	0.000000
0.000000				
50%	0.000000	0.000000	0.000000	0.000000
0.000000				
75%	0.000000	0.000000	0.000000	0.000000
0.000000				
max	1.000000	1.000000	1.000000	1.000000
1.000000				

	Age	Gender	Subscription_Length_Months
Monthly_Bill \			
count	80000.000000	80000.000000	80000.000000
80000.000000			
mean	44.016225	0.497762	12.48990
65.076958			
std	15.278733	0.499998	6.91766
20.227098			
min	18.000000	0.000000	1.00000
30.000000			
25%	31.000000	0.000000	6.00000
47.600000			
50%	44.000000	0.000000	12.00000
65.030000			
75%	57.000000	1.000000	18.00000
82.700000			
max	70.000000	1.000000	24.00000
100.000000			

	Total_Usage_GB
count	80000.000000
mean	274.662787
std	130.510754
min	50.000000
25%	161.000000
50%	274.000000
75%	388.000000
max	500.000000

*# Feature scaling using maximum values*

```

X_train['Age'] = X_train['Age']/70
X_train['Subscription_Length_Months'] =
X_train['Subscription_Length_Months']/24
X_train['Monthly_Bill'] = X_train['Monthly_Bill']/100
X_train['Total_Usage_GB'] = X_train['Total_Usage_GB']/500
X_train

```

	Chicago	Houston	Los Angeles	Miami	New York	Age
Gender \						

75220	0.0	0.0	0.0	0.0	1.0	0.771429
0						
48955	0.0	0.0	0.0	0.0	1.0	0.400000
1						
44966	1.0	0.0	0.0	0.0	0.0	0.814286
1						
13568	0.0	1.0	0.0	0.0	0.0	0.271429
1						
92727	0.0	0.0	0.0	1.0	0.0	0.800000
0						
...	...	...	...	...	...	...
.						
6265	0.0	0.0	0.0	1.0	0.0	0.500000
1						
54886	1.0	0.0	0.0	0.0	0.0	0.800000
1						
76820	0.0	1.0	0.0	0.0	0.0	0.985714
1						
860	1.0	0.0	0.0	0.0	0.0	0.785714
1						
15795	0.0	0.0	1.0	0.0	0.0	0.371429
0						
	Subscription_Length_Months		Monthly_Bill		Total_Usage_GB	
75220	0.208333		0.8450		0.410	
48955	1.000000		0.8206		0.478	
44966	0.500000		0.5229		0.124	
13568	0.791667		0.3257		0.346	
92727	0.333333		0.3352		0.628	
...	...		...		...	
6265	0.875000		0.6733		0.470	
54886	0.541667		0.8540		0.694	
76820	0.083333		0.7624		0.642	
860	0.500000		0.8919		0.630	
15795	0.708333		0.7041		0.670	
[80000 rows x 10 columns]						

# Model Building

## Logistic Regression

```
print(X_train)
print(X_train.shape)
```

	Chicago	Houston	Los Angeles	Miami	New York	Age
Gender \						

75220	0.0	0.0	0.0	0.0	1.0	0.771429
0						
48955	0.0	0.0	0.0	0.0	1.0	0.400000
1						
44966	1.0	0.0	0.0	0.0	0.0	0.814286
1						
13568	0.0	1.0	0.0	0.0	0.0	0.271429
1						
92727	0.0	0.0	0.0	1.0	0.0	0.800000
0						
...	...	...	...	...	...	...
.						
6265	0.0	0.0	0.0	1.0	0.0	0.500000
1						
54886	1.0	0.0	0.0	0.0	0.0	0.800000
1						
76820	0.0	1.0	0.0	0.0	0.0	0.985714
1						
860	1.0	0.0	0.0	0.0	0.0	0.785714
1						
15795	0.0	0.0	1.0	0.0	0.0	0.371429
0						

	Subscription_Length_Months	Monthly_Bill	Total_Usage_GB
75220	0.208333	0.8450	0.410
48955	1.000000	0.8206	0.478
44966	0.500000	0.5229	0.124
13568	0.791667	0.3257	0.346
92727	0.333333	0.3352	0.628
...	...	...	...
6265	0.875000	0.6733	0.470
54886	0.541667	0.8540	0.694
76820	0.083333	0.7624	0.642
860	0.500000	0.8919	0.630
15795	0.708333	0.7041	0.670

```
[80000 rows x 10 columns]
(80000, 10)
```

```
y_train.shape
(80000,)
```

```
lg_model = LogisticRegression()
lg_model = lg_model.fit(X_train,y_train)
```

```
y_pred_lg_train = lg_model.predict(X_train)
print(classification_report(y_train,y_pred_lg_train))
print(y_train)
print(y_pred_lg_train)
```

	precision	recall	f1-score	support
0	0.50	0.62	0.56	40142
1	0.50	0.38	0.44	39858
accuracy			0.50	80000
macro avg	0.50	0.50	0.50	80000
weighted avg	0.50	0.50	0.50	80000
[1 1 1 ... 1 1 0]				
[1 1 1 ... 0 0 0]				

Total overall accuracy is 50%

Precision for True Negatives is 51% and that of True Positives is 51%

The performance is average

## Support Vector Machines

```
svm_model = SVC()
svm_model = svm_model.fit(X_train,y_train)

y_pred_svm = svm_model.predict(X_test)
print(classification_report(y_test,y_pred_svm))
```

The performance of SVM is quite similar to Logistic Regression

However in some cases it performs better, like in case of recall and f1-score

## Random Forest Classifier

```
rand_forest_model = RandomForestClassifier()
rand_forest_model.fit(X_train,y_train)

y_pred_forest = rand_forest_model.predict(X_test)
print(classification_report(y_test,y_pred_forest))

rand_forest_model.score(X_test,y_test)
```

Overall the Random Forest is performing worse than both Logistic Regression and SVM with an accuracy of less than 50%

## Fine Tuning Random Forest

Checking the performance by increasing the number of Trees

```
rand_forest_model = RandomForestClassifier(n_estimators=20)
rand_forest_model.fit(X_train,y_train)

rand_forest_model.score(X_test,y_test)
```

For 20 trees the model does not improve that much

Further increasing the number of Trees

```
rand_forest_model = RandomForestClassifier(n_estimators=40)
rand_forest_model.fit(X_train,y_train)

rand_forest_model.score(X_test,y_test)

rand_forest_pred = rand_forest_model.predict(X_test)
```


Best performance occurs with  $n\_estimators = 40$

Further increasing the value leads to diminishing returns

```
# Plotting Confusion Matrix for Random Forest
cm = confusion_matrix(y_test,rand_forest_pred)

cm

# Visualizing the confusion matrix
%matplotlib inline
import seaborn as sns
sns.heatmap(cm,annot=True)
plt.xlabel('Predicted')
plt.ylabel('Ground Truth')
plt.show()
```



- 5200 True Negatives were correctly predicted
- 4700 True Positives were correctly predicted
- A similar number of predicted values were wrongly predicted
- This implies the accuracy of the model is around 50% as seen previously

## Neural Network

```
def plot_loss(history):  
    plt.plot(history.history['loss'], label='loss')  
    plt.plot(history.history['val_loss'], label='val_loss')  
    plt.xlabel('Epoch')  
    plt.ylabel('Binary Crossentropy')  
    plt.legend()  
    plt.grid(True)  
    plt.show()  
  
def plot_accuracy(history):  
    plt.plot(history.history['accuracy'], label='accuracy')  
    plt.plot(history.history['val_accuracy'], label='val_accuracy')
```



```

plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)
plt.show()

nn_model = tf.keras.Sequential([
    tf.keras.layers.Dense(64,activation='relu'),
    tf.keras.layers.Dense(32,activation='relu'),
    tf.keras.layers.Dense(1,activation='sigmoid'),
])

nn_model.compile(optimizer=tf.keras.optimizers.Adam(0.001),loss='binary_crossentropy',metrics=['accuracy'])

history = nn_model.fit(
    X_train, y_train,epochs=100,batch_size=32,validation_split=0.2
)

Epoch 1/100
2000/2000 [=====] - 7s 3ms/step - loss:
0.6945 - accuracy: 0.5004 - val_loss: 0.6940 - val_accuracy: 0.5013
Epoch 2/100
2000/2000 [=====] - 6s 3ms/step - loss:
0.6934 - accuracy: 0.5054 - val_loss: 0.6933 - val_accuracy: 0.5003
Epoch 3/100
2000/2000 [=====] - 5s 3ms/step - loss:
0.6932 - accuracy: 0.5054 - val_loss: 0.6933 - val_accuracy: 0.4963
Epoch 4/100
2000/2000 [=====] - 7s 4ms/step - loss:
0.6931 - accuracy: 0.5064 - val_loss: 0.6938 - val_accuracy: 0.4970
Epoch 5/100
2000/2000 [=====] - 14s 7ms/step - loss:
0.6931 - accuracy: 0.5089 - val_loss: 0.6933 - val_accuracy: 0.5019
Epoch 6/100
2000/2000 [=====] - 5s 2ms/step - loss:
0.6930 - accuracy: 0.5090 - val_loss: 0.6934 - val_accuracy: 0.5002
Epoch 7/100
2000/2000 [=====] - 5s 3ms/step - loss:
0.6930 - accuracy: 0.5102 - val_loss: 0.6935 - val_accuracy: 0.5017
Epoch 8/100
2000/2000 [=====] - 6s 3ms/step - loss:
0.6929 - accuracy: 0.5088 - val_loss: 0.6938 - val_accuracy: 0.4980
Epoch 9/100
2000/2000 [=====] - 5s 2ms/step - loss:
0.6928 - accuracy: 0.5119 - val_loss: 0.6938 - val_accuracy: 0.4986
Epoch 10/100
2000/2000 [=====] - 6s 3ms/step - loss:
0.6928 - accuracy: 0.5132 - val_loss: 0.6934 - val_accuracy: 0.5038

```

Epoch 11/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6927 - accuracy: 0.5126 - val\_loss: 0.6933 - val\_accuracy: 0.5006  
Epoch 12/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6926 - accuracy: 0.5121 - val\_loss: 0.6935 - val\_accuracy: 0.4992  
Epoch 13/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6925 - accuracy: 0.5164 - val\_loss: 0.6933 - val\_accuracy: 0.5008  
Epoch 14/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6925 - accuracy: 0.5138 - val\_loss: 0.6940 - val\_accuracy: 0.5011  
Epoch 15/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6925 - accuracy: 0.5160 - val\_loss: 0.6939 - val\_accuracy: 0.4972  
Epoch 16/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6924 - accuracy: 0.5158 - val\_loss: 0.6946 - val\_accuracy: 0.5008  
Epoch 17/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6923 - accuracy: 0.5186 - val\_loss: 0.6944 - val\_accuracy: 0.4979  
Epoch 18/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6922 - accuracy: 0.5168 - val\_loss: 0.6937 - val\_accuracy: 0.5001  
Epoch 19/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6922 - accuracy: 0.5178 - val\_loss: 0.6952 - val\_accuracy: 0.5012  
Epoch 20/100  
2000/2000 [=====] - 12s 6ms/step - loss:  
0.6921 - accuracy: 0.5194 - val\_loss: 0.6942 - val\_accuracy: 0.4980  
Epoch 21/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6921 - accuracy: 0.5193 - val\_loss: 0.6946 - val\_accuracy: 0.4947  
Epoch 22/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6920 - accuracy: 0.5195 - val\_loss: 0.6949 - val\_accuracy: 0.4964  
Epoch 23/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6919 - accuracy: 0.5204 - val\_loss: 0.6952 - val\_accuracy: 0.4964  
Epoch 24/100  
2000/2000 [=====] - 4s 2ms/step - loss:  
0.6919 - accuracy: 0.5203 - val\_loss: 0.6942 - val\_accuracy: 0.5058  
Epoch 25/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6919 - accuracy: 0.5213 - val\_loss: 0.6946 - val\_accuracy: 0.4999  
Epoch 26/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6918 - accuracy: 0.5200 - val\_loss: 0.6944 - val\_accuracy: 0.5020  
Epoch 27/100

2000/2000 [=====] - 6s 3ms/step - loss:  
0.6917 - accuracy: 0.5214 - val\_loss: 0.6947 - val\_accuracy: 0.4955  
Epoch 28/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6918 - accuracy: 0.5209 - val\_loss: 0.6946 - val\_accuracy: 0.5010  
Epoch 29/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6916 - accuracy: 0.5235 - val\_loss: 0.6942 - val\_accuracy: 0.5050  
Epoch 30/100  
2000/2000 [=====] - 4s 2ms/step - loss:  
0.6915 - accuracy: 0.5211 - val\_loss: 0.6944 - val\_accuracy: 0.5026  
Epoch 31/100  
2000/2000 [=====] - 7s 3ms/step - loss:  
0.6915 - accuracy: 0.5241 - val\_loss: 0.6949 - val\_accuracy: 0.5011  
Epoch 32/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6915 - accuracy: 0.5228 - val\_loss: 0.6949 - val\_accuracy: 0.4990  
Epoch 33/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6913 - accuracy: 0.5247 - val\_loss: 0.6949 - val\_accuracy: 0.5014  
Epoch 34/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6912 - accuracy: 0.5262 - val\_loss: 0.6954 - val\_accuracy: 0.5024  
Epoch 35/100  
2000/2000 [=====] - 4s 2ms/step - loss:  
0.6913 - accuracy: 0.5248 - val\_loss: 0.6951 - val\_accuracy: 0.5016  
Epoch 36/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6913 - accuracy: 0.5249 - val\_loss: 0.6952 - val\_accuracy: 0.5029  
Epoch 37/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6913 - accuracy: 0.5258 - val\_loss: 0.6946 - val\_accuracy: 0.5073  
Epoch 38/100  
2000/2000 [=====] - 7s 3ms/step - loss:  
0.6911 - accuracy: 0.5249 - val\_loss: 0.6950 - val\_accuracy: 0.5039  
Epoch 39/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6911 - accuracy: 0.5239 - val\_loss: 0.6943 - val\_accuracy: 0.5076  
Epoch 40/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6910 - accuracy: 0.5266 - val\_loss: 0.6950 - val\_accuracy: 0.5056  
Epoch 41/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6909 - accuracy: 0.5260 - val\_loss: 0.6948 - val\_accuracy: 0.5027  
Epoch 42/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6910 - accuracy: 0.5267 - val\_loss: 0.6949 - val\_accuracy: 0.5063  
Epoch 43/100  
2000/2000 [=====] - 5s 2ms/step - loss:

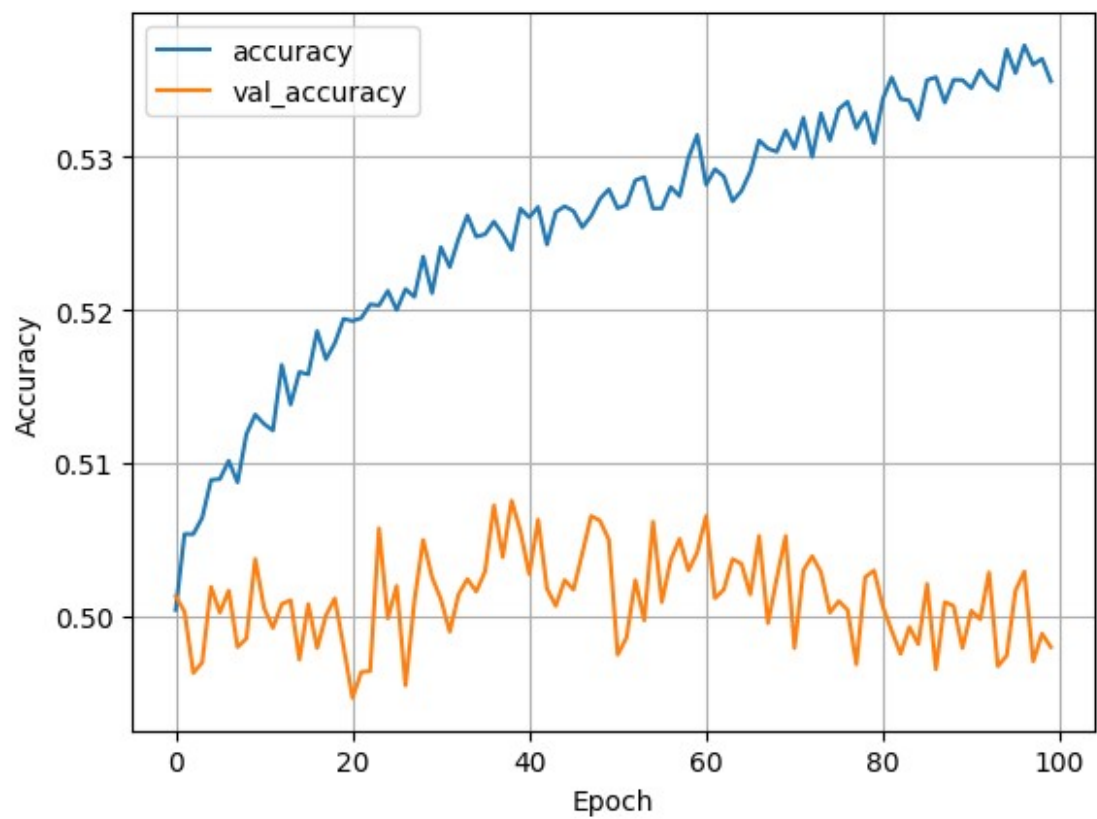

0.6909 - accuracy: 0.5243 - val\_loss: 0.6953 - val\_accuracy: 0.5018  
Epoch 44/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6906 - accuracy: 0.5264 - val\_loss: 0.6960 - val\_accuracy: 0.5007  
Epoch 45/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6907 - accuracy: 0.5268 - val\_loss: 0.6955 - val\_accuracy: 0.5024  
Epoch 46/100  
2000/2000 [=====] - 4s 2ms/step - loss:  
0.6906 - accuracy: 0.5265 - val\_loss: 0.6965 - val\_accuracy: 0.5017  
Epoch 47/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6907 - accuracy: 0.5254 - val\_loss: 0.6954 - val\_accuracy: 0.5041  
Epoch 48/100  
2000/2000 [=====] - 4s 2ms/step - loss:  
0.6906 - accuracy: 0.5261 - val\_loss: 0.6960 - val\_accuracy: 0.5066  
Epoch 49/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6905 - accuracy: 0.5272 - val\_loss: 0.6951 - val\_accuracy: 0.5063  
Epoch 50/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6904 - accuracy: 0.5279 - val\_loss: 0.6956 - val\_accuracy: 0.5050  
Epoch 51/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6904 - accuracy: 0.5266 - val\_loss: 0.6961 - val\_accuracy: 0.4975  
Epoch 52/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6904 - accuracy: 0.5269 - val\_loss: 0.6965 - val\_accuracy: 0.4986  
Epoch 53/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6903 - accuracy: 0.5285 - val\_loss: 0.6962 - val\_accuracy: 0.5024  
Epoch 54/100  
2000/2000 [=====] - 4s 2ms/step - loss:  
0.6902 - accuracy: 0.5287 - val\_loss: 0.6954 - val\_accuracy: 0.4997  
Epoch 55/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6902 - accuracy: 0.5266 - val\_loss: 0.6959 - val\_accuracy: 0.5062  
Epoch 56/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6903 - accuracy: 0.5266 - val\_loss: 0.6964 - val\_accuracy: 0.5009  
Epoch 57/100  
2000/2000 [=====] - 4s 2ms/step - loss:  
0.6901 - accuracy: 0.5280 - val\_loss: 0.6965 - val\_accuracy: 0.5037  
Epoch 58/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6901 - accuracy: 0.5274 - val\_loss: 0.6969 - val\_accuracy: 0.5051  
Epoch 59/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6901 - accuracy: 0.5299 - val\_loss: 0.6960 - val\_accuracy: 0.5030

Epoch 60/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6900 - accuracy: 0.5314 - val\_loss: 0.6961 - val\_accuracy: 0.5042  
Epoch 61/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6900 - accuracy: 0.5282 - val\_loss: 0.6953 - val\_accuracy: 0.5066  
Epoch 62/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6899 - accuracy: 0.5292 - val\_loss: 0.6966 - val\_accuracy: 0.5012  
Epoch 63/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6899 - accuracy: 0.5287 - val\_loss: 0.6966 - val\_accuracy: 0.5017  
Epoch 64/100  
2000/2000 [=====] - 7s 3ms/step - loss:  
0.6899 - accuracy: 0.5271 - val\_loss: 0.6971 - val\_accuracy: 0.5038  
Epoch 65/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6897 - accuracy: 0.5278 - val\_loss: 0.6967 - val\_accuracy: 0.5034  
Epoch 66/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6897 - accuracy: 0.5290 - val\_loss: 0.6960 - val\_accuracy: 0.5014  
Epoch 67/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6896 - accuracy: 0.5311 - val\_loss: 0.6968 - val\_accuracy: 0.5052  
Epoch 68/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6896 - accuracy: 0.5305 - val\_loss: 0.6970 - val\_accuracy: 0.4996  
Epoch 69/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6897 - accuracy: 0.5303 - val\_loss: 0.6968 - val\_accuracy: 0.5025  
Epoch 70/100  
2000/2000 [=====] - 4s 2ms/step - loss:  
0.6894 - accuracy: 0.5317 - val\_loss: 0.6972 - val\_accuracy: 0.5052  
Epoch 71/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6895 - accuracy: 0.5305 - val\_loss: 0.6972 - val\_accuracy: 0.4979  
Epoch 72/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6893 - accuracy: 0.5325 - val\_loss: 0.6969 - val\_accuracy: 0.5030  
Epoch 73/100  
2000/2000 [=====] - 4s 2ms/step - loss:  
0.6894 - accuracy: 0.5300 - val\_loss: 0.6970 - val\_accuracy: 0.5039  
Epoch 74/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6892 - accuracy: 0.5328 - val\_loss: 0.6972 - val\_accuracy: 0.5029  
Epoch 75/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6893 - accuracy: 0.5311 - val\_loss: 0.6971 - val\_accuracy: 0.5002  
Epoch 76/100

2000/2000 [=====] - 4s 2ms/step - loss:  
0.6892 - accuracy: 0.5331 - val\_loss: 0.6966 - val\_accuracy: 0.5010  
Epoch 77/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6891 - accuracy: 0.5336 - val\_loss: 0.6978 - val\_accuracy: 0.5004  
Epoch 78/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6892 - accuracy: 0.5319 - val\_loss: 0.6979 - val\_accuracy: 0.4969  
Epoch 79/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6891 - accuracy: 0.5329 - val\_loss: 0.6981 - val\_accuracy: 0.5026  
Epoch 80/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6891 - accuracy: 0.5309 - val\_loss: 0.6976 - val\_accuracy: 0.5030  
Epoch 81/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6889 - accuracy: 0.5338 - val\_loss: 0.6979 - val\_accuracy: 0.5006  
Epoch 82/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6889 - accuracy: 0.5352 - val\_loss: 0.6980 - val\_accuracy: 0.4991  
Epoch 83/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6889 - accuracy: 0.5338 - val\_loss: 0.6972 - val\_accuracy: 0.4976  
Epoch 84/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6887 - accuracy: 0.5337 - val\_loss: 0.6984 - val\_accuracy: 0.4993  
Epoch 85/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6888 - accuracy: 0.5324 - val\_loss: 0.6980 - val\_accuracy: 0.4982  
Epoch 86/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6888 - accuracy: 0.5350 - val\_loss: 0.6980 - val\_accuracy: 0.5021  
Epoch 87/100  
2000/2000 [=====] - 4s 2ms/step - loss:  
0.6886 - accuracy: 0.5352 - val\_loss: 0.6986 - val\_accuracy: 0.4966  
Epoch 88/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6885 - accuracy: 0.5335 - val\_loss: 0.6992 - val\_accuracy: 0.5009  
Epoch 89/100  
2000/2000 [=====] - 7s 4ms/step - loss:  
0.6885 - accuracy: 0.5350 - val\_loss: 0.6987 - val\_accuracy: 0.5007  
Epoch 90/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6886 - accuracy: 0.5350 - val\_loss: 0.7000 - val\_accuracy: 0.4979  
Epoch 91/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6885 - accuracy: 0.5345 - val\_loss: 0.6989 - val\_accuracy: 0.5004  
Epoch 92/100  
2000/2000 [=====] - 5s 3ms/step - loss:

```
0.6884 - accuracy: 0.5356 - val_loss: 0.6994 - val_accuracy: 0.4998
Epoch 93/100
2000/2000 [=====] - 4s 2ms/step - loss:
0.6885 - accuracy: 0.5348 - val_loss: 0.6987 - val_accuracy: 0.5029
Epoch 94/100
2000/2000 [=====] - 5s 3ms/step - loss:
0.6883 - accuracy: 0.5344 - val_loss: 0.6989 - val_accuracy: 0.4967
Epoch 95/100
2000/2000 [=====] - 5s 3ms/step - loss:
0.6884 - accuracy: 0.5370 - val_loss: 0.6987 - val_accuracy: 0.4974
Epoch 96/100
2000/2000 [=====] - 4s 2ms/step - loss:
0.6880 - accuracy: 0.5355 - val_loss: 0.6982 - val_accuracy: 0.5017
Epoch 97/100
2000/2000 [=====] - 6s 3ms/step - loss:
0.6883 - accuracy: 0.5373 - val_loss: 0.6992 - val_accuracy: 0.5029
Epoch 98/100
2000/2000 [=====] - 5s 3ms/step - loss:
0.6882 - accuracy: 0.5360 - val_loss: 0.6996 - val_accuracy: 0.4971
Epoch 99/100
2000/2000 [=====] - 5s 2ms/step - loss:
0.6881 - accuracy: 0.5364 - val_loss: 0.6988 - val_accuracy: 0.4989
Epoch 100/100
2000/2000 [=====] - 6s 3ms/step - loss:
0.6881 - accuracy: 0.5349 - val_loss: 0.6988 - val_accuracy: 0.4980

plot_loss(history)
plot_accuracy(history)
```





The Neural Networks are performing well as the loss is clearly decreasing and the accuracy is increasing with each iteration.

The accuracy however is 53% which is better than the other algorithms

## Fine Tuning model parameters

testing on learning rate = 0.005 and 64 neurons in the 2nd layer

```
nn_model = tf.keras.Sequential([
    tf.keras.layers.Dense(64,activation='relu'),
    tf.keras.layers.Dense(64,activation='relu'),
    tf.keras.layers.Dense(1,activation='sigmoid'),
])

nn_model.compile(optimizer=tf.keras.optimizers.Adam(0.001),loss='binary_crossentropy',metrics=['accuracy'])

history = nn_model.fit(
    X_train, y_train,epochs=100,batch_size=32,validation_split=0.2
)
```

```
Epoch 1/100
2000/2000 [=====] - 9s 3ms/step - loss:
0.6943 - accuracy: 0.4991 - val_loss: 0.6945 - val_accuracy: 0.4971
Epoch 2/100
2000/2000 [=====] - 5s 3ms/step - loss:
0.6935 - accuracy: 0.5016 - val_loss: 0.6935 - val_accuracy: 0.4930
Epoch 3/100
2000/2000 [=====] - 7s 3ms/step - loss:
0.6932 - accuracy: 0.5040 - val_loss: 0.6932 - val_accuracy: 0.5009
Epoch 4/100
2000/2000 [=====] - 5s 3ms/step - loss:
0.6932 - accuracy: 0.5063 - val_loss: 0.6933 - val_accuracy: 0.4960
Epoch 5/100
2000/2000 [=====] - 6s 3ms/step - loss:
0.6930 - accuracy: 0.5095 - val_loss: 0.6933 - val_accuracy: 0.4999
Epoch 6/100
2000/2000 [=====] - 4s 2ms/step - loss:
0.6929 - accuracy: 0.5086 - val_loss: 0.6935 - val_accuracy: 0.4974
Epoch 7/100
2000/2000 [=====] - 5s 3ms/step - loss:
```

0.6929 - accuracy: 0.5092 - val\_loss: 0.6935 - val\_accuracy: 0.5039  
Epoch 8/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6928 - accuracy: 0.5127 - val\_loss: 0.6940 - val\_accuracy: 0.5004  
Epoch 9/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6928 - accuracy: 0.5110 - val\_loss: 0.6935 - val\_accuracy: 0.4998  
Epoch 10/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6927 - accuracy: 0.5121 - val\_loss: 0.6937 - val\_accuracy: 0.4978  
Epoch 11/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6926 - accuracy: 0.5139 - val\_loss: 0.6940 - val\_accuracy: 0.4959  
Epoch 12/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6927 - accuracy: 0.5147 - val\_loss: 0.6940 - val\_accuracy: 0.4956  
Epoch 13/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6925 - accuracy: 0.5145 - val\_loss: 0.6942 - val\_accuracy: 0.4996  
Epoch 14/100  
2000/2000 [=====] - 4s 2ms/step - loss:  
0.6925 - accuracy: 0.5160 - val\_loss: 0.6943 - val\_accuracy: 0.5024  
Epoch 15/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6924 - accuracy: 0.5142 - val\_loss: 0.6941 - val\_accuracy: 0.5026  
Epoch 16/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6923 - accuracy: 0.5175 - val\_loss: 0.6941 - val\_accuracy: 0.4992  
Epoch 17/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6923 - accuracy: 0.5194 - val\_loss: 0.6938 - val\_accuracy: 0.5008  
Epoch 18/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6923 - accuracy: 0.5155 - val\_loss: 0.6941 - val\_accuracy: 0.5031  
Epoch 19/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6922 - accuracy: 0.5163 - val\_loss: 0.6944 - val\_accuracy: 0.4983  
Epoch 20/100  
2000/2000 [=====] - 4s 2ms/step - loss:  
0.6920 - accuracy: 0.5194 - val\_loss: 0.6945 - val\_accuracy: 0.4989  
Epoch 21/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6920 - accuracy: 0.5197 - val\_loss: 0.6946 - val\_accuracy: 0.5006  
Epoch 22/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6920 - accuracy: 0.5200 - val\_loss: 0.6943 - val\_accuracy: 0.5029  
Epoch 23/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6919 - accuracy: 0.5200 - val\_loss: 0.6947 - val\_accuracy: 0.4979

Epoch 24/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6918 - accuracy: 0.5190 - val\_loss: 0.6942 - val\_accuracy: 0.5044  
Epoch 25/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6918 - accuracy: 0.5214 - val\_loss: 0.6943 - val\_accuracy: 0.5033  
Epoch 26/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6919 - accuracy: 0.5200 - val\_loss: 0.6946 - val\_accuracy: 0.5023  
Epoch 27/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6918 - accuracy: 0.5213 - val\_loss: 0.6948 - val\_accuracy: 0.4994  
Epoch 28/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6916 - accuracy: 0.5231 - val\_loss: 0.6946 - val\_accuracy: 0.4991  
Epoch 29/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6916 - accuracy: 0.5228 - val\_loss: 0.6949 - val\_accuracy: 0.5017  
Epoch 30/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6915 - accuracy: 0.5220 - val\_loss: 0.6948 - val\_accuracy: 0.5042  
Epoch 31/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6914 - accuracy: 0.5218 - val\_loss: 0.6946 - val\_accuracy: 0.5031  
Epoch 32/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6912 - accuracy: 0.5248 - val\_loss: 0.6950 - val\_accuracy: 0.5019  
Epoch 33/100  
2000/2000 [=====] - 7s 4ms/step - loss:  
0.6913 - accuracy: 0.5228 - val\_loss: 0.6949 - val\_accuracy: 0.4988  
Epoch 34/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6912 - accuracy: 0.5228 - val\_loss: 0.6948 - val\_accuracy: 0.5034  
Epoch 35/100  
2000/2000 [=====] - 7s 3ms/step - loss:  
0.6911 - accuracy: 0.5256 - val\_loss: 0.6953 - val\_accuracy: 0.5042  
Epoch 36/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6910 - accuracy: 0.5268 - val\_loss: 0.6956 - val\_accuracy: 0.5005  
Epoch 37/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6909 - accuracy: 0.5254 - val\_loss: 0.6955 - val\_accuracy: 0.5017  
Epoch 38/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6910 - accuracy: 0.5242 - val\_loss: 0.6953 - val\_accuracy: 0.4998  
Epoch 39/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6910 - accuracy: 0.5256 - val\_loss: 0.6954 - val\_accuracy: 0.4982  
Epoch 40/100

2000/2000 [=====] - 6s 3ms/step - loss:  
0.6908 - accuracy: 0.5276 - val\_loss: 0.6949 - val\_accuracy: 0.5038  
Epoch 41/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6908 - accuracy: 0.5265 - val\_loss: 0.6955 - val\_accuracy: 0.5021  
Epoch 42/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6906 - accuracy: 0.5276 - val\_loss: 0.6953 - val\_accuracy: 0.5021  
Epoch 43/100  
2000/2000 [=====] - 7s 3ms/step - loss:  
0.6907 - accuracy: 0.5263 - val\_loss: 0.6957 - val\_accuracy: 0.5026  
Epoch 44/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6905 - accuracy: 0.5281 - val\_loss: 0.6960 - val\_accuracy: 0.5027  
Epoch 45/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6906 - accuracy: 0.5262 - val\_loss: 0.6954 - val\_accuracy: 0.5029  
Epoch 46/100  
2000/2000 [=====] - 7s 3ms/step - loss:  
0.6905 - accuracy: 0.5260 - val\_loss: 0.6964 - val\_accuracy: 0.5014  
Epoch 47/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6904 - accuracy: 0.5277 - val\_loss: 0.6957 - val\_accuracy: 0.5038  
Epoch 48/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6903 - accuracy: 0.5282 - val\_loss: 0.6966 - val\_accuracy: 0.5045  
Epoch 49/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6902 - accuracy: 0.5294 - val\_loss: 0.6959 - val\_accuracy: 0.4991  
Epoch 50/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6901 - accuracy: 0.5310 - val\_loss: 0.6959 - val\_accuracy: 0.5013  
Epoch 51/100  
2000/2000 [=====] - 7s 3ms/step - loss:  
0.6902 - accuracy: 0.5293 - val\_loss: 0.6956 - val\_accuracy: 0.5034  
Epoch 52/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6901 - accuracy: 0.5285 - val\_loss: 0.6959 - val\_accuracy: 0.4988  
Epoch 53/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6900 - accuracy: 0.5301 - val\_loss: 0.6972 - val\_accuracy: 0.5008  
Epoch 54/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6900 - accuracy: 0.5315 - val\_loss: 0.6961 - val\_accuracy: 0.5019  
Epoch 55/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6899 - accuracy: 0.5302 - val\_loss: 0.6962 - val\_accuracy: 0.5059  
Epoch 56/100  
2000/2000 [=====] - 5s 3ms/step - loss:

0.6898 - accuracy: 0.5313 - val\_loss: 0.6976 - val\_accuracy: 0.4976  
Epoch 57/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6897 - accuracy: 0.5291 - val\_loss: 0.6968 - val\_accuracy: 0.4998  
Epoch 58/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6898 - accuracy: 0.5316 - val\_loss: 0.6962 - val\_accuracy: 0.5038  
Epoch 59/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6898 - accuracy: 0.5317 - val\_loss: 0.6970 - val\_accuracy: 0.5018  
Epoch 60/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6897 - accuracy: 0.5314 - val\_loss: 0.6963 - val\_accuracy: 0.5092  
Epoch 61/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6897 - accuracy: 0.5298 - val\_loss: 0.6960 - val\_accuracy: 0.5077  
Epoch 62/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6897 - accuracy: 0.5318 - val\_loss: 0.6966 - val\_accuracy: 0.5060  
Epoch 63/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6896 - accuracy: 0.5298 - val\_loss: 0.6966 - val\_accuracy: 0.5039  
Epoch 64/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6894 - accuracy: 0.5328 - val\_loss: 0.6967 - val\_accuracy: 0.5016  
Epoch 65/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6895 - accuracy: 0.5319 - val\_loss: 0.6969 - val\_accuracy: 0.5030  
Epoch 66/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6894 - accuracy: 0.5307 - val\_loss: 0.6967 - val\_accuracy: 0.5077  
Epoch 67/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6892 - accuracy: 0.5322 - val\_loss: 0.6969 - val\_accuracy: 0.5069  
Epoch 68/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6891 - accuracy: 0.5351 - val\_loss: 0.6978 - val\_accuracy: 0.4979  
Epoch 69/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6892 - accuracy: 0.5332 - val\_loss: 0.6970 - val\_accuracy: 0.5065  
Epoch 70/100  
2000/2000 [=====] - 7s 3ms/step - loss:  
0.6892 - accuracy: 0.5332 - val\_loss: 0.6969 - val\_accuracy: 0.5017  
Epoch 71/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6890 - accuracy: 0.5342 - val\_loss: 0.6974 - val\_accuracy: 0.5001  
Epoch 72/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6891 - accuracy: 0.5349 - val\_loss: 0.6969 - val\_accuracy: 0.5054

Epoch 73/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6889 - accuracy: 0.5341 - val\_loss: 0.6973 - val\_accuracy: 0.5059  
Epoch 74/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6889 - accuracy: 0.5346 - val\_loss: 0.6985 - val\_accuracy: 0.5049  
Epoch 75/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6888 - accuracy: 0.5364 - val\_loss: 0.6975 - val\_accuracy: 0.5027  
Epoch 76/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6889 - accuracy: 0.5352 - val\_loss: 0.6980 - val\_accuracy: 0.5021  
Epoch 77/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6888 - accuracy: 0.5357 - val\_loss: 0.6967 - val\_accuracy: 0.5069  
Epoch 78/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6887 - accuracy: 0.5360 - val\_loss: 0.6977 - val\_accuracy: 0.5030  
Epoch 79/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6888 - accuracy: 0.5348 - val\_loss: 0.6973 - val\_accuracy: 0.5042  
Epoch 80/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6886 - accuracy: 0.5342 - val\_loss: 0.6969 - val\_accuracy: 0.5044  
Epoch 81/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6885 - accuracy: 0.5353 - val\_loss: 0.6970 - val\_accuracy: 0.5052  
Epoch 82/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6885 - accuracy: 0.5357 - val\_loss: 0.6972 - val\_accuracy: 0.4997  
Epoch 83/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6885 - accuracy: 0.5357 - val\_loss: 0.6984 - val\_accuracy: 0.4984  
Epoch 84/100  
2000/2000 [=====] - 6s 3ms/step - loss:  
0.6885 - accuracy: 0.5379 - val\_loss: 0.6974 - val\_accuracy: 0.5038  
Epoch 85/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6883 - accuracy: 0.5365 - val\_loss: 0.6983 - val\_accuracy: 0.5039  
Epoch 86/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6884 - accuracy: 0.5351 - val\_loss: 0.6971 - val\_accuracy: 0.5055  
Epoch 87/100  
2000/2000 [=====] - 5s 3ms/step - loss:  
0.6883 - accuracy: 0.5364 - val\_loss: 0.6979 - val\_accuracy: 0.5043  
Epoch 88/100  
2000/2000 [=====] - 5s 2ms/step - loss:  
0.6883 - accuracy: 0.5376 - val\_loss: 0.6984 - val\_accuracy: 0.5032  
Epoch 89/100

```

2000/2000 [=====] - 6s 3ms/step - loss:
0.6882 - accuracy: 0.5357 - val_loss: 0.6984 - val_accuracy: 0.5043
Epoch 90/100
2000/2000 [=====] - 5s 2ms/step - loss:
0.6882 - accuracy: 0.5364 - val_loss: 0.6974 - val_accuracy: 0.5026
Epoch 91/100
2000/2000 [=====] - 6s 3ms/step - loss:
0.6881 - accuracy: 0.5393 - val_loss: 0.6975 - val_accuracy: 0.5063
Epoch 92/100
2000/2000 [=====] - 7s 4ms/step - loss:
0.6882 - accuracy: 0.5373 - val_loss: 0.6976 - val_accuracy: 0.5081
Epoch 93/100
2000/2000 [=====] - 6s 3ms/step - loss:
0.6879 - accuracy: 0.5377 - val_loss: 0.6981 - val_accuracy: 0.5036
Epoch 94/100
2000/2000 [=====] - 6s 3ms/step - loss:
0.6880 - accuracy: 0.5372 - val_loss: 0.6977 - val_accuracy: 0.5050
Epoch 95/100
2000/2000 [=====] - 7s 3ms/step - loss:
0.6880 - accuracy: 0.5378 - val_loss: 0.6974 - val_accuracy: 0.5080
Epoch 96/100
2000/2000 [=====] - 5s 3ms/step - loss:
0.6881 - accuracy: 0.5363 - val_loss: 0.6982 - val_accuracy: 0.5042
Epoch 97/100
2000/2000 [=====] - 7s 3ms/step - loss:
0.6879 - accuracy: 0.5391 - val_loss: 0.6982 - val_accuracy: 0.5005
Epoch 98/100
2000/2000 [=====] - 5s 3ms/step - loss:
0.6879 - accuracy: 0.5382 - val_loss: 0.6978 - val_accuracy: 0.5023
Epoch 99/100
2000/2000 [=====] - 6s 3ms/step - loss:
0.6877 - accuracy: 0.5392 - val_loss: 0.6985 - val_accuracy: 0.5024
Epoch 100/100
2000/2000 [=====] - 5s 2ms/step - loss:
0.6878 - accuracy: 0.5392 - val_loss: 0.6987 - val_accuracy: 0.5039

```

Accuracy is approx 53% which is better than the rest of the algorithms

Script to find the best possible parameters for neural network

```

def plot_history(history):
    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(10, 4))
    ax1.plot(history.history['loss'], label='loss')
    ax1.plot(history.history['val_loss'], label='val_loss')
    ax1.set_xlabel('Epoch')
    ax1.set_ylabel('Binary crossentropy')
    ax1.grid(True)

```

```

ax2.plot(history.history['accuracy'], label='accuracy')
ax2.plot(history.history['val_accuracy'], label='val_accuracy')
ax2.set_xlabel('Epoch')
ax2.set_ylabel('Accuracy')
ax2.grid(True)

plt.show()

def train_model(X_train, y_train, num_nodes, dropout_prob, lr,
batch_size, epochs):
    nn_model = tf.keras.Sequential([
        tf.keras.layers.Dense(num_nodes, activation='relu'),
        tf.keras.layers.Dropout(dropout_prob),
        tf.keras.layers.Dense(num_nodes, activation='relu'),
        tf.keras.layers.Dropout(dropout_prob),
        tf.keras.layers.Dense(1, activation='sigmoid')
    ])

    nn_model.compile(optimizer=tf.keras.optimizers.Adam(lr),
loss='binary_crossentropy',
                    metrics=['accuracy'])
    history = nn_model.fit(
        X_train, y_train, epochs=epochs, batch_size=batch_size,
validation_split=0.2, verbose=0
    )

    return nn_model, history

least_val_loss = float('inf')
least_loss_model = None
epochs=100


"""
    UNCOMMENT THE CODE BELOW TO RUN
"""

# for num_nodes in [16, 32, 64]:
#     for dropout_prob in [0, 0.2]:
#         for lr in [0.01, 0.005, 0.001]:
#             for batch_size in [32, 64]:
#                 print(f"{num_nodes} nodes, dropout {dropout_prob}, lr {lr},
batch size {batch_size}")
#                 model, history = train_model(X_train, y_train, num_nodes,
dropout_prob, lr, batch_size, epochs)
#                 plot_history(history)
#                 val_loss = model.evaluate(X_test, y_test)[0]
#                 if val_loss < least_val_loss:
#                     least_val_loss = val_loss
#                     least_loss_model = model


```





16 nodes, dropout 0, lr 0.01, batch size 32




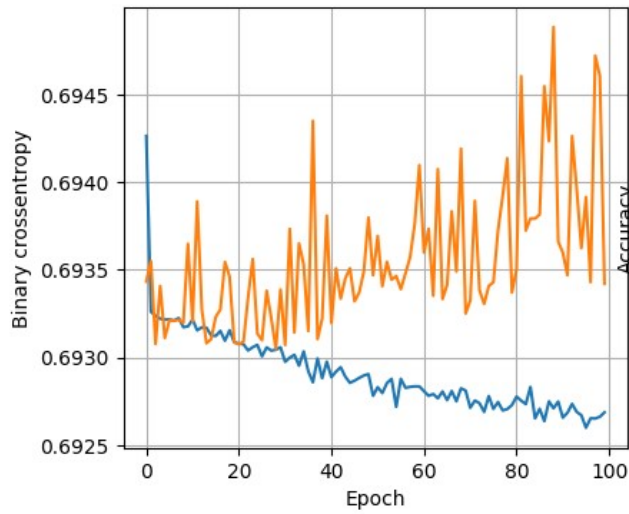
625/625 [=====] - 1s 2ms/step - loss: 0.6932  
- accuracy: 0.4961  
16 nodes, dropout 0, lr 0.01, batch size 64




625/625 [=====] - 1s 2ms/step - loss: 0.6932  
- accuracy: 0.5039  
16 nodes, dropout 0, lr 0.005, batch size 32




625/625 [=====] - 1s 2ms/step - loss: 0.6931  
 - accuracy: 0.5039  
 16 nodes, dropout 0, lr 0.005, batch size 64





625/625 [=====] - 1s 2ms/step - loss: 0.6937  
 - accuracy: 0.4975  
 16 nodes, dropout 0, lr 0.001, batch size 32





625/625 [=====] - 1s 2ms/step - loss: 0.6958  
 - accuracy: 0.4988  
 16 nodes, dropout 0, lr 0.001, batch size 64



625/625 [=====] - 1s 2ms/step - loss: 0.6947  
 - accuracy: 0.4988  
 16 nodes, dropout 0.2, lr 0.01, batch size 32





625/625 [=====] - 1s 2ms/step - loss: 0.6933  
 - accuracy: 0.5039  
 16 nodes, dropout 0.2, lr 0.01, batch size 64





625/625 [=====] - 3s 2ms/step - loss: 0.6933  
 - accuracy: 0.4961  
 16 nodes, dropout 0.2, lr 0.005, batch size 32









625/625 [=====] - 1s 2ms/step - loss: 0.6933  
 - accuracy: 0.5039  
 16 nodes, dropout 0.2, lr 0.005, batch size 64



625/625 [=====] - 2s 3ms/step - loss: 0.6931  
 - accuracy: 0.5039  
 16 nodes, dropout 0.2, lr 0.001, batch size 32



```
625/625 [=====] - 1s 2ms/step - loss: 0.6933
- accuracy: 0.5016
16 nodes, dropout 0.2, lr 0.001, batch size 64
```



```
625/625 [=====] - 1s 2ms/step - loss: 0.6937
- accuracy: 0.4982
32 nodes, dropout 0, lr 0.01, batch size 32
```

```
-----
-----
KeyboardInterrupt                                Traceback (most recent call
last)
<ipython-input-96-c7802061c1b1> in <cell line: 4>()
      7         for batch_size in [32, 64]:
      8             print(f"{num_nodes} nodes, dropout {dropout_prob}, lr
{lr}, batch size {batch_size}")
```

```

----> 9         model, history = train_model(X_train, y_train,
num_nodes, dropout_prob, lr, batch_size, epochs)
      10         plot_history(history)
      11         val_loss = model.evaluate(X_test, y_test)[0]

<ipython-input-94-97d214b36a92> in train_model(X_train, y_train,
num_nodes, dropout_prob, lr, batch_size, epochs)
      10     nn_model.compile(optimizer=tf.keras.optimizers.Adam(lr),
loss='binary_crossentropy',
      11                         metrics=['accuracy'])
--> 12     history = nn_model.fit(
      13         X_train, y_train, epochs=epochs, batch_size=batch_size,
validation_split=0.2, verbose=0
      14     )

/usr/local/lib/python3.10/dist-packages/keras/utils/traceback_utils.py
in error_handler(*args, **kwargs)
      63         filtered_tb = None
      64         try:
--> 65             return fn(*args, **kwargs)
      66         except Exception as e:
      67             filtered_tb =
_process_traceback_frames(e.__traceback__)

/usr/local/lib/python3.10/dist-packages/keras/engine/training.py in
fit(self, x, y, batch_size, epochs, verbose, callbacks,
validation_split, validation_data, shuffle, class_weight,
sample_weight, initial_epoch, steps_per_epoch, validation_steps,
validation_batch_size, validation_freq, max_queue_size, workers,
use_multiprocessing)
    1683         ):
    1684
callbacks.on_train_batch_begin(step)
-> 1685         tmp_logs =
self.train_function(iterator)
    1686         if data_handler.should_sync:
    1687             context.async_wait()

/usr/local/lib/python3.10/dist-packages/tensorflow/python/util/traceba
ck_utils.py in error_handler(*args, **kwargs)
    148         filtered_tb = None
    149         try:
--> 150             return fn(*args, **kwargs)
    151         except Exception as e:
    152             filtered_tb = _process_traceback_frames(e.__traceback__)

/usr/local/lib/python3.10/dist-packages/tensorflow/python/eager/polymo
rphic_function/polymorphic_function.py in __call__(self, *args,
**kws)
    892

```

```

893         with OptionalXlaContext(self._jit_compile):
--> 894             result = self._call(*args, **kwds)
895
896         new_tracing_count =
self.experimental_get_tracing_count()

/usr/local/lib/python3.10/dist-packages/tensorflow/python/eager/polymo
rphic_function/polymorphic_function.py in _call(self, *args, **kwds)
924         # In this case we have created variables on the first
call, so we run the
925         # defunned version which is guaranteed to never create
variables.
--> 926         return self._no_variable_creation_fn(*args, **kwds) #
pylint: disable=not-callable
927     elif self._variable_creation_fn is not None:
928         # Release the lock early so that multiple threads can
perform the call

/usr/local/lib/python3.10/dist-packages/tensorflow/python/eager/polymo
rphic_function/tracing_compiler.py in __call__(self, *args, **kwargs)
141         (concrete_function,
142          filtered_flat_args) = self._maybe_define_function(args,
kwargs)
--> 143         return concrete_function._call_flat(
144             filtered_flat_args,
captured_inputs=concrete_function.captured_inputs) # pylint:
disable=protected-access
145

/usr/local/lib/python3.10/dist-packages/tensorflow/python/eager/polymo
rphic_function/monomorphic_function.py in _call_flat(self, args,
captured_inputs, cancellation_manager)
1755         and executing_eagerly):
1756         # No tape is watching; skip to running the function.
-> 1757         return
self._build_call_outputs(self._inference_function.call(
1758             ctx, args,
cancellation_manager=cancellation_manager))
1759         forward_backward =
self._select_forward_and_backward_functions(

/usr/local/lib/python3.10/dist-packages/tensorflow/python/eager/polymo
rphic_function/monomorphic_function.py in call(self, ctx, args,
cancellation_manager)
379         with _InterpolateFunctionError(self):
380             if cancellation_manager is None:
--> 381                 outputs = execute.execute(
382                     str(self.signature.name),
383                     num_outputs=self._num_outputs,

```



```
/usr/local/lib/python3.10/dist-packages/tensorflow/python/eager/execute.py in quick_execute(op_name, num_outputs, inputs, attrs, ctx, name)
    50     try:
    51         ctx.ensure_initialized()
--> 52         tensors = pywrap_tfe.TFE_Py_Execute(ctx._handle,
device_name, op_name,
    53                                     inputs, attrs,
num_outputs)
    54     except core._NotOkStatusException as e:
```

KeyboardInterrupt:

```
import pickle
pickle.dump(lg_model, open('lg_model.pkl', 'wb'))

loaded_model = pickle.load(open('lg_model.pkl', 'rb'))
```